

# face-emotion-recognition

March 31, 2024

## 0.1 Face Emotion Recognition

```
[10]: #importing important libraries
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import os

# Importing Deep Learning Libraries

from keras.preprocessing.image import load_img, img_to_array
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import
    ↪Dense, Input, Dropout, GlobalAveragePooling2D, Flatten, Conv2D, BatchNormalization, Activation, Max
from keras.models import Model, Sequential
from tensorflow.keras.optimizers import Adam, SGD, RMSprop
```

## 1 Displaying Images

```
[11]: #taking inputs images for training
picture_size = 48
folder_path = "../input/face-expression-recognition-dataset/images/"

[12]: #checking disgust image
expression = 'disgust'

plt.figure(figsize= (12,12))
for i in range(1, 10, 1):
    plt.subplot(3,3,i)
    img = load_img(folder_path+"train/"+expression+"/"+
                    os.listdir(folder_path + "train/" + expression)[i],
    ↪target_size=(picture_size, picture_size))
    plt.imshow(img)
plt.show()
```



## 2 Making Training and Validation Data

```
[13]: #splitting data into train, test and validation set
batch_size = 128

datagen_train = ImageDataGenerator()
datagen_val = ImageDataGenerator()

train_set = datagen_train.flow_from_directory(folder_path+"train",
                                              target_size = (
        ↳(picture_size,picture_size),
                                              color_mode = "grayscale",
```

```

                                batch_size=batch_size,
                                class_mode='categorical',
                                shuffle=True)

test_set = datagen_val.flow_from_directory(folder_path+"validation",
                                target_size = (
↪(picture_size,picture_size),

                                color_mode = "grayscale",
                                batch_size=batch_size,
                                class_mode='categorical',
                                shuffle=False)

```

Found 28821 images belonging to 7 classes.

Found 7066 images belonging to 7 classes.

### 3 Model Building

- Model = sequential : A Sequential model is appropriate for a plain stack of layers where each layer has exactly one input tensor and one output tensor.
- Activation = relu :The rectified linear activation function or ReLU for short is a piecewise linear function that will output the input directly if it is positive, otherwise, it will output zero. It has become the default activation function for many types of neural networks because a model that uses it is easier to train and often achieves better performance.
- padding = The padding parameter of the Keras Conv2D class can take one of two values: 'valid' or 'same'. Setting the value to "valid" parameter means that the input volume is not zero-padded and the spatial dimensions are allowed to reduce via the natural application of convolution.
- Maxpooling = Maximum pooling, or max pooling, is a pooling operation that calculates the maximum, or largest, value in each patch of each feature map. The results are down sampled or pooled feature maps that highlight the most present feature in the patch, not the average presence of the feature in the case of average pooling.
- Batch normalization = Batch normalization is a technique for training very deep neural networks that standardizes the inputs to a layer for each mini-batch. This has the effect of stabilizing the learning process and dramatically reducing the number of training epochs required to train deep networks.
- Dropout = Dropout is a technique used to prevent a model from overfitting. Dropout works by randomly setting the outgoing edges of hidden units (neurons that make up hidden layers) to 0 at each update of the training phase.
- Adam = Adam can be looked at as a combination of RMSprop and Stochastic Gradient Descent with momentum. It uses the squared gradients to scale the learning rate like RMSprop and it takes advantage of momentum by using moving average of the gradient instead of gradient itself like SGD with momentum.

- SGD = Stochastic Gradient Descent (SGD) addresses both of these issues by following the negative gradient of the objective after seeing only a single or a few training examples. The use of SGD In the neural network setting is motivated by the high cost of running back propagation over the full training set
- RMSprop = RMSprop is a gradient based optimization technique used in training neural networks. ... This normalization balances the step size (momentum), decreasing the step for large gradients to avoid exploding, and increasing the step for small gradients to avoid vanishing.

```
[14]: #building model with 7 classes
from tensorflow.keras.optimizers import Adam,SGD,RMSprop

no_of_classes = 7

model = Sequential()

#1st CNN layer
model.add(Conv2D(64,(3,3),padding = 'same',input_shape = (48,48,1)))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

#2nd CNN layer
model.add(Conv2D(128,(5,5),padding = 'same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout (0.25))

#3rd CNN layer
model.add(Conv2D(512,(3,3),padding = 'same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout (0.25))

#4th CNN layer
model.add(Conv2D(512,(3,3), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
```

```

#Fully connected 1st layer
model.add(Dense(256))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

# Fully connected layer 2nd layer
model.add(Dense(512))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

model.add(Dense(no_of_classes, activation='softmax'))

opt = Adam(lr = 0.0001)
model.compile(optimizer=opt,loss='categorical_crossentropy',
              metrics=['accuracy'])
model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 48, 48, 64)	640
batch_normalization_6 (Batch Normalization)	(None, 48, 48, 64)	256
activation_6 (Activation)	(None, 48, 48, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 24, 24, 64)	0
dropout_6 (Dropout)	(None, 24, 24, 64)	0
conv2d_5 (Conv2D)	(None, 24, 24, 128)	204928
batch_normalization_7 (Batch Normalization)	(None, 24, 24, 128)	512
activation_7 (Activation)	(None, 24, 24, 128)	0
max_pooling2d_5 (MaxPooling2D)	(None, 12, 12, 128)	0
dropout_7 (Dropout)	(None, 12, 12, 128)	0
conv2d_6 (Conv2D)	(None, 12, 12, 512)	590336

-----		
batch_normalization_8 (Batch Normalization)	(None, 12, 12, 512)	2048
-----		
activation_8 (Activation)	(None, 12, 12, 512)	0
-----		
max_pooling2d_6 (MaxPooling2D)	(None, 6, 6, 512)	0
-----		
dropout_8 (Dropout)	(None, 6, 6, 512)	0
-----		
conv2d_7 (Conv2D)	(None, 6, 6, 512)	2359808
-----		
batch_normalization_9 (Batch Normalization)	(None, 6, 6, 512)	2048
-----		
activation_9 (Activation)	(None, 6, 6, 512)	0
-----		
max_pooling2d_7 (MaxPooling2D)	(None, 3, 3, 512)	0
-----		
dropout_9 (Dropout)	(None, 3, 3, 512)	0
-----		
flatten_1 (Flatten)	(None, 4608)	0
-----		
dense_3 (Dense)	(None, 256)	1179904
-----		
batch_normalization_10 (Batch Normalization)	(None, 256)	1024
-----		
activation_10 (Activation)	(None, 256)	0
-----		
dropout_10 (Dropout)	(None, 256)	0
-----		
dense_4 (Dense)	(None, 512)	131584
-----		
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
-----		
activation_11 (Activation)	(None, 512)	0
-----		
dropout_11 (Dropout)	(None, 512)	0
-----		
dense_5 (Dense)	(None, 7)	3591
=====		

Total params: 4,478,727

Trainable params: 4,474,759

Non-trainable params: 3,968

```

/opt/conda/lib/python3.7/site-packages/keras/optimizer_v2/optimizer_v2.py:356:
UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  "The `lr` argument is deprecated, use `learning_rate` instead.")

```

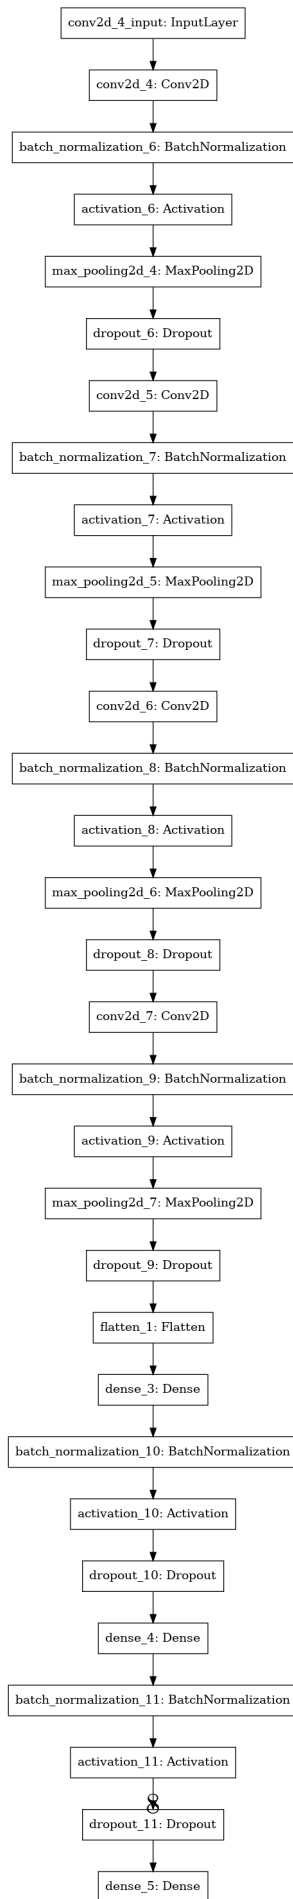
## 4 Visualize model

The `plot_model()` function in Keras will create a plot of your network. This function takes a few useful arguments:

- `model`: (required) The model that you wish to plot.
- `to_file`: (required) The name of the file to which to save the plot.
- `show_shapes`: (optional, defaults to `False`) Whether or not to show the output shapes of each layer.
- `show_layer_names`: (optional, defaults to `True`) Whether or not to show the name for each layer.

```
[15]: #visualizing the model
import tensorflow as tf
tf.keras.utils.plot_model(
    model,
    to_file="model.png",
    show_shapes=False,
    show_dtype=False,
    show_layer_names=True,
    rankdir="TB",
    expand_nested=False,
    dpi=96,
    layer_range=None,
)
```

[15]:





## 5 Fitting the Model with Training and Validation Data

```
[16]: #importing tensorflow library and package
from tensorflow.keras.optimizers import RMSprop,SGD,Adam
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping,ReduceLROnPlateau

checkpoint = ModelCheckpoint("./model.h5", monitor='val_acc', verbose=1,save_best_only=True, mode='max')
#Stopping training when a monitored metric has stopped improving.
early_stopping = EarlyStopping(monitor='val_loss',
                                min_delta=0,
                                patience=3,
                                verbose=1,
                                restore_best_weights=True
                                )

reduce_learningrate = ReduceLROnPlateau(monitor='val_loss',
                                          factor=0.2,
                                          patience=3,
                                          verbose=1,
                                          min_delta=0.0001)

callbacks_list = [early_stopping,checkpoint,reduce_learningrate]

epochs = 50

model.compile(loss='categorical_crossentropy',
              optimizer = Adam(lr=0.001),
              metrics=['accuracy'])
```

/opt/conda/lib/python3.7/site-packages/keras/optimizer\_v2/optimizer\_v2.py:356:  
UserWarning: The `lr` argument is deprecated, use `learning\_rate` instead.  
"The `lr` argument is deprecated, use `learning\_rate` instead.")

```
[17]: #fitting model with 48 epoch
history = model.fit_generator(generator=train_set,
                              steps_per_epoch=train_set.n//train_set.
                              batch_size,

                              epochs=epochs,
                              validation_data = test_set,
                              validation_steps = test_set.n//test_set.
                              batch_size,
```

```
callbacks=callbacks_list
)
```

```
/opt/conda/lib/python3.7/site-packages/keras/engine/training.py:1972:
UserWarning: `Model.fit_generator` is deprecated and will be removed in a future
version. Please use `Model.fit`, which supports generators.
```

```
warnings.warn("`Model.fit_generator` is deprecated and "
```

```
Epoch 1/50
```

```
225/225 [=====] - 32s 136ms/step - loss: 1.7579 -
accuracy: 0.3280 - val_loss: 1.7229 - val_accuracy: 0.3463
```

```
Epoch 2/50
```

```
225/225 [=====] - 30s 134ms/step - loss: 1.4099 -
accuracy: 0.4605 - val_loss: 1.4641 - val_accuracy: 0.4230
```

```
Epoch 3/50
```

```
225/225 [=====] - 30s 131ms/step - loss: 1.2680 -
accuracy: 0.5133 - val_loss: 1.4188 - val_accuracy: 0.4530
```

```
Epoch 4/50
```

```
225/225 [=====] - 35s 154ms/step - loss: 1.1786 -
accuracy: 0.5528 - val_loss: 1.2958 - val_accuracy: 0.4989
```

```
Epoch 5/50
```

```
225/225 [=====] - 30s 134ms/step - loss: 1.1190 -
accuracy: 0.5758 - val_loss: 1.4361 - val_accuracy: 0.4652
```

```
Epoch 6/50
```

```
225/225 [=====] - 30s 131ms/step - loss: 1.0743 -
accuracy: 0.5957 - val_loss: 1.1260 - val_accuracy: 0.5714
```

```
Epoch 7/50
```

```
225/225 [=====] - 30s 135ms/step - loss: 1.0183 -
accuracy: 0.6164 - val_loss: 1.1967 - val_accuracy: 0.5480
```

```
Epoch 8/50
```

```
225/225 [=====] - 30s 133ms/step - loss: 0.9873 -
accuracy: 0.6271 - val_loss: 1.1446 - val_accuracy: 0.5628
```

```
Epoch 9/50
```

```
225/225 [=====] - 30s 132ms/step - loss: 0.9402 -
accuracy: 0.6472 - val_loss: 1.0745 - val_accuracy: 0.5977
```

```
Epoch 10/50
```

```
225/225 [=====] - 30s 133ms/step - loss: 0.8981 -
accuracy: 0.6634 - val_loss: 1.0688 - val_accuracy: 0.6026
```

```
Epoch 11/50
```

```
225/225 [=====] - 30s 132ms/step - loss: 0.8555 -
accuracy: 0.6794 - val_loss: 1.0726 - val_accuracy: 0.6010
```

```
Epoch 12/50
```

```
225/225 [=====] - 30s 133ms/step - loss: 0.8154 -
accuracy: 0.6949 - val_loss: 1.0688 - val_accuracy: 0.6036
```

```
Epoch 13/50
```

```
225/225 [=====] - 29s 130ms/step - loss: 0.7711 -
accuracy: 0.7118 - val_loss: 1.1160 - val_accuracy: 0.6021
```

Epoch 00013: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.  
Epoch 14/50  
225/225 [=====] - 30s 135ms/step - loss: 0.6395 -  
accuracy: 0.7639 - val\_loss: 0.9925 - val\_accuracy: 0.6514  
Epoch 15/50  
225/225 [=====] - 30s 133ms/step - loss: 0.5905 -  
accuracy: 0.7821 - val\_loss: 1.0203 - val\_accuracy: 0.6575  
Epoch 16/50  
225/225 [=====] - 29s 130ms/step - loss: 0.5570 -  
accuracy: 0.7954 - val\_loss: 1.0347 - val\_accuracy: 0.6477  
Epoch 17/50  
225/225 [=====] - 30s 132ms/step - loss: 0.5260 -  
accuracy: 0.8070 - val\_loss: 1.0348 - val\_accuracy: 0.6511  
Restoring model weights from the end of the best epoch.


Epoch 00017: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.  
Epoch 00017: early stopping

## 6 Plotting Accuracy & Loss

```
[18]: #plotting graph to check accuracy and loss
plt.style.use('dark_background')

plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : Adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend(loc='lower right')
plt.show()
```



[ ]: